AI-Powered Legal Document Analysis and Compliance Assistant

Dharin Dave[1], Harshraj Chouhan[1], Sakshi Jambhulkar[1], Yash Agarwal[1] and Santwana Gudadhe[1]

[1] Pimpri Chinchwad College of Engineering, Pune, India
lncs@springer.com

Abstract. An AI-powered approach for automating legal document interpretation and improving compliance management in intricate regulatory contexts is presented in this study. The suggested system makes effective use of a vector database and a fine-tuned large language model (LLM) to store, retrieve, and analyze legal documents semantically. The platform's orchestration layer, which is created using cutting-edge frameworks like LangChain, is at its core and manages duties like contract review, risk assessment, and entity extraction and summarization using natural language processing. In order to reduce human review efforts and mitigate legal risks, the system also includes a generative AI module for proactive contract evaluation and a real-time compliance monitoring dashboard that provides actionable insights. The chosen model performs better in context relevance, groundedness, and response relevance, according to comparative performance evaluations performed across several language models (including Llama 3.1, GPT-4, Gemini 1.5, and Mistral 12B). The outcomes demonstrate the usefulness of the suggested method in offering trustworthy and effective legal document analysis, as well as its potential for domain-specific flexibility and scalable, cloud-based implementation.

Keywords: Compliance, LLM, LangChain, Vector Database, NLP, Summarization

1 Introduction

The legal sector is now dealing with a fast-changing regulatory environment on top of an ever-increasing amount and complexity of legal papers. Conventional legal analysis techniques, which mostly rely on manual review, are showing themselves to be ineffective and prone to mistakes, which frequently leads to overlooked compliance concerns and increased risk exposure [3][4]. There has never been a more pressing demand for automated, dependable, and scalable solutions, especially for businesses that must handle complex legal requirements in several countries [11][15].

Legal document analysis might be revolutionized by recent developments in artificial intelligence, especially in the areas of natural language processing and large-scale language models [2][13][16][22]. To ensure high standards of accuracy and relevance while streamlining the review process, this paper presents an AI-Powered Legal Document Analysis and Compliance Assistant. With the help of cutting-edge big language models and a vector database for effective document embedding and retrieval, the system can extract important legal provisions, track compliance levels, and identify any contractual hazards [12][20][21]. The system's ability to adjust to domain-specific legal difficulties is further enhanced by the addition of an orchestration layer, which makes

it easier to integrate different activities including contract review, risk assessment, and document summarizing [1][5][10].

The suggested system incorporates a generative AI module for proactive contract review and a real-time compliance monitoring dashboard that offers actionable insights, in addition to its sturdy design [3][6][14]. Tests against other models, such as GPT-4, Gemini 1.5, and Mistral 12B, have repeatedly shown that the chosen method is superior when it comes to important performance metrics like response relevance, context relevance, and groundedness [7][8][9][19]. These empirical results highlight the platform's potential to revolutionize compliance monitoring and legal document management procedures in a scalable and economical way [18][23].

This is how the rest of the paper is organized. Section 2 examines relevant research on the use of AI and deep learning for compliance monitoring and legal document interpretation. The suggested technique is described in Section 3 and is divided into three main modules: real-time compliance tracking, generative contract assessment, and automated legal document processing. The experimental design, evaluation measures, and outcomes of our performance evaluations are shown in Section 4. The study is finally concluded in Section 5, which also discusses possible future research and development possibilities and summarizes the contributions of this work

2 Literature Survey

The field of legal artificial intelligence (LegalAI) is concerned with the use of AI, namely natural language processing (NLP), in the legal field. By automating repetitive procedures, this seeks to lessen the workload for attorneys and other legal professionals [4].

[4] explains the potential advantages of NLP in the legal field based on past advancements, present uses, and potential future advances.

Integrating legal information with Graph Neural Networks (GNNs) is discussed in [5]. The difficulties posed by the vast volumes of data stored in legal documents are addressed in [6]. The authors have put forth a novel strategy that uses explainable AI tools to improve conventional processes. These methods have demonstrated an increase in the effectiveness and openness of the legal document review process [6].

[3] discusses how big data analytics and legal systems interact. This state-of-the-art research provides a summary of studies on the future scope of contemporary legal and big data thought. It includes the issues in big data analytics within the legal system and emphasizes the importance of quality and value in legal services and best practices for ensuring the quality of legal documents.

A summary of the incorporation of large language models in the legal field can be found in [2]. It primarily focuses on the uses, difficulties, and potential future applications of LLMs in legal situations. The concerns of data privacy, bias, interpretability, and regulation are discussed in this study. Taking care of these issues will aid in the creation of effective LLMs in the legal field [2].

The capabilities of pre-trained language models to generalize across various legal domains and transfer knowledge between datasets annotated with different type systems are discussed [1]. Predicting the functions of sentences in court rulings is the primary focus of this study; this is a laborious, expensive, and time-consuming process that must be done by hand. Large language models have been refined on several legal datasets because of the need to handle this. According to the authors, this method may address the cold-start issue in interactive or active learning environments [1].

3 Proposed Methodology

3.1 Extraction Module

The Extraction module is a robust and versatile Python tool designed for extracting and processing content from PDF documents. It utilizes the pdfplumber library for reading PDFs, pandas for handling tabular data, and re for searching and processing text patterns. Additionally, it integrates Python's built-in logging module to provide structured error handling and debugging capabilities. This class is particularly valuable in scenarios where automated extraction of structured data from PDFs is required, such as legal document processing, financial reporting, research paper analysis, and compliance checks.

One of the core functionalities of the extraction module is its ability to extract text while preserving formatting. Unlike simple text extraction methods that return unstructured output, this implementation reconstructs text by considering word positioning and maintaining document flow. It processes multiple pages, ensuring that extracted content retains coherence. Furthermore, it marks page breaks to distinguish sections clearly, making it especially useful for multi-page documents where structured formatting is essential.

Beyond text extraction, the class also supports the extraction of tables, converting them into structured pandas data frames. Many PDFs contain tabular data, which can be difficult to parse accurately. By leveraging pdfplumber library, the class identifies tables, cleans the extracted data, and structures it in a way that is easily accessible for further analysis. This functionality is particularly useful for processing financial reports, invoices, and research data, where tables are a primary source of information.

The class also includes a metadata extraction feature, which retrieves key document properties such as the title, author, creation date, and other metadata fields. This can be helpful in organizing and indexing large collections of PDFs, as well as in verifying the authenticity of documents. Metadata extraction is widely used in applications like document classification, information retrieval, and digital forensics, where understanding a document's origins is critical.

Another important feature of the extraction module is its ability to search for text patterns within a document using regular expressions. This allows users to perform advanced text searches, such as identifying specific keywords, extracting structured data or scanning for sensitive information. The class also provides contextual information for each match, including the page number, position in text, and surrounding words.

This makes it especially useful for legal discovery, compliance audits, and automated content analysis.

3.2 Summarization Module

The summarization module follows a structured workflow designed to systematically load, process, and summarize a given PDF document. The process begins when the user provides the file path of the document along with an optional API key for Groq's language model. If the API key is not explicitly supplied, the system retrieves it from a secure environment file to ensure sensitive credentials are managed safely without being hardcoded into the script.

Once initialized, the document is loaded into memory using a specialized PDF processing tool that extracts text from the file. The extracted content is stored within the system, making it readily available for further processing. This approach assumes that the entire document can be handled in one go. However, if the document is particularly large, additional techniques may be required to divide it into smaller segments to prevent exceeding token limitations when interacting with the language model.

Before generating the summary, the language model needs to be set up with specific configurations. A high-performance model is selected, with settings that ensure deterministic and factual responses. This setup is particularly useful for legal document summarization, where accuracy and consistency are critical. The model is initialized only once during the session, optimizing performance and reducing redundant computations.

With both the document and the language model ready, the summarization process can be initiated. The system first ensures that all necessary components are properly set up before proceeding. A structured prompt is then crafted, instructing the language model to simplify complex legal language into clear, layman-friendly terms. This ensures that even individuals without legal expertise can understand the document's content. The structured nature of the prompt helps maintain clarity and conciseness in the generated summary.

To process the document, a dedicated pipeline is constructed that seamlessly integrates the extracted text with the language model. The system then analyses the content and produces a summary that is designed to be clear, concise, and free from unnecessary legal jargon. Once generated, the summary is returned as a text output that can be displayed, stored, or further processed as needed.

By following this structured approach, the summarization module effectively converts complex legal documents into simplified summaries, making them more accessible to a wider audience. While the current implementation provides a solid foundation, further refinements—such as improved document segmentation, enhanced retrieval mechanisms, and support for multiple models—could enhance its efficiency and scalability, especially when dealing with large legal texts.

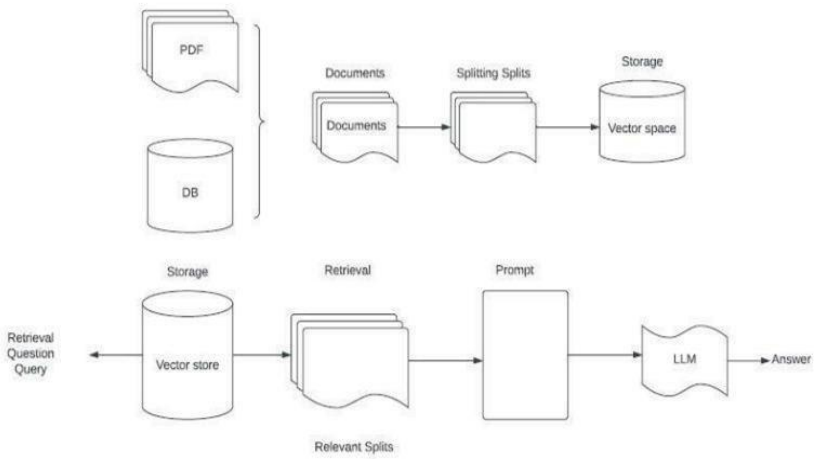

Fig. 1. Overview of Proposed Methodology

3.3 Querying Module

Fig. 1 follows a structured process to enable efficient retrieval and question-answering for legal texts. The system is designed to load, process, and store legal documents in a way that allows users to query them for relevant information. The process begins when the user provides the file path of a legal document, along with an optional API key for Groq's language model. If the API key is not explicitly supplied, the system securely loads it from an environment file to maintain confidentiality and avoid hardcoding sensitive credentials.

Once initialized, the system processes the document by first extracting text from the PDF file. The text is then divided into smaller chunks using a text-splitting mechanism to ensure that the information can be effectively stored and retrieved. The splitting process maintains some overlap between chunks to preserve context, which is particularly useful when answering detailed legal queries.

After the text is split, it is embedded into a vector representation using a deterministic embedding model. These embeddings are stored in a FAISS vector database, which allows for efficient similarity-based retrieval of relevant document sections when a query is made. This setup ensures that instead of searching through the entire document manually, the system can quickly find the most relevant sections based on the user's query.

The system then initializes a large language model (LLM) from Groq, configured with a temperature setting of 0 to ensure factual and deterministic responses. The LLM is used to generate answers based on retrieved document segments, maintaining accuracy and reliability in legal contexts.

To facilitate effective querying, a Retrieval-Augmented Generation (RAG) pipeline is set up. When the user submits a question, the system first retrieves the most relevant

document sections using similarity search. These retrieved sections are formatted and passed into a structured prompt that instructs the language model on how to respond. The prompt is carefully designed to ensure that the model only provides answers based on the given context, avoids making up information, and remains professional and clear in its responses. If the query is unrelated to the document, the system explicitly states that the question is not relevant.

With this setup, Fig. 2 enables users to efficiently extract relevant information from lengthy legal texts. Instead of reading through complex legal documents manually, users can ask specific questions and receive precise, context-aware answers. While the current implementation provides a strong foundation, further enhancements—such as integrating more advanced embeddings, handling multiple documents, and improving retrieval precision—could further enhance the system's effectiveness in legal and compliance-related use cases.

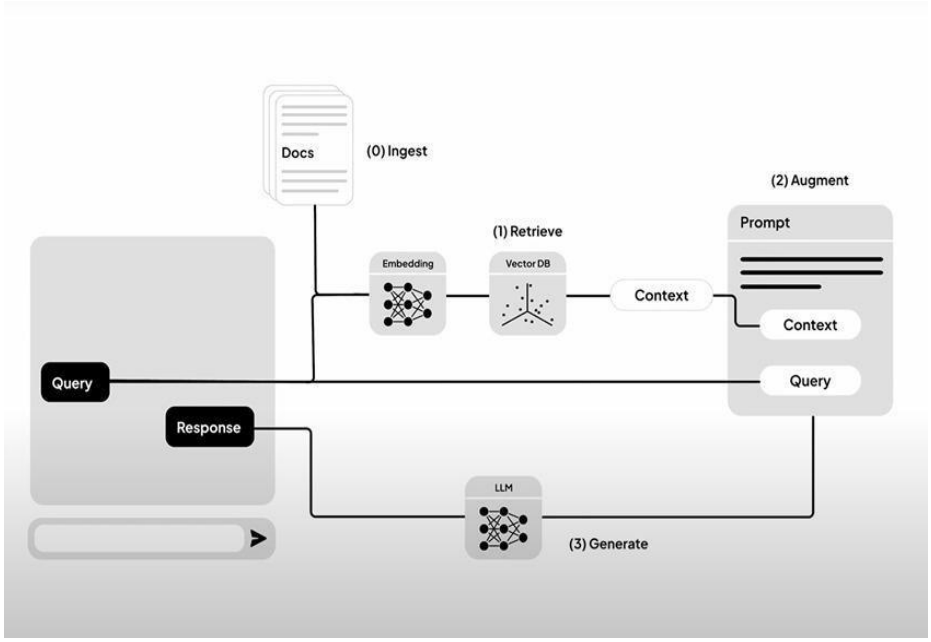

Fig. 2. Proposed Methodology for Querying Module

4 Experiments and Results

4.1 Experimental Design

To evaluate the efficacy of our AI-powered legal document analysis and compliance assistant, we conducted a comprehensive assessment of four state-of-the-art language models: Llama 3.1, GPT-4, Gemini 1.5, and Mistral 12B. The evaluation framework was designed to measure performance across three critical dimensions essential for

legal and compliance applications: context relevance, groundedness, and answer relevance. These metrics were selected to address the unique challenges inherent in legal document processing, where precision, factual accuracy, and contextual understanding are paramount.

To evaluate the efficacy of our AI-powered legal document analysis and compliance assistant, we conducted a comprehensive assessment of four state-of-the-art language models: Llama 3.1, GPT-4, Gemini 1.5, and Mistral 12B. The evaluation framework was designed to measure performance across three critical dimensions essential for legal and compliance applications: context relevance, groundedness, and answer relevance. These metrics were selected to address the unique challenges inherent in legal document processing, where precision, factual accuracy, and contextual understanding are paramount.

4.2 Performance Metrics

Context relevance quantifies a model's ability to comprehend and respond appropriately to the nuanced contextual elements of legal inquiries. This metric is particularly vital in the legal domain, where misinterpretation of context can lead to significant compliance risks. Our evaluation revealed that Llama 3.1 achieved superior performance with a score of 0.98, demonstrating exceptional capability in contextual understanding. GPT-4 matched this score, while Gemini 1.5 and Mistral 12B scored 0.92 and 0.90 respectively, indicating a notable performance gap.

Groundedness measures the extent to which model responses remain anchored to verified information without unwarranted extrapolation or hallucination. In legal and compliance contexts, where strict adherence to regulatory facts is critical, Llama 3.1 demonstrated superior performance with a score of 0.97. This represents a significant advantage over GPT-4 (0.94), Gemini 1.5 (0.93), and Mistral 12B (0.91). The superior groundedness of Llama 3.1 substantiates its reliability for compliance-sensitive applications where factual accuracy cannot be compromised.

Answer relevance evaluates how directly and effectively a model addresses the specific query posed. This metric is essential for assessing the practical utility of responses in legal contexts. Our analysis indicates that Llama 3.1 outperformed all competing models with an exceptional score of 0.99, compared to GPT-4 (0.96), Gemini 1.5 (0.95), and Mistral 12B (0.91). This suggests that Llama 3.1 provides more focused and pertinent responses to legal and compliance inquiries.

4.3 Comparative Analysis

The comparative performance of the evaluated models reveals several significant insights. Llama 3.1 (Fig. 3) consistently demonstrated superior capabilities across all three metrics, achieving an average score of 0.98, which surpasses the next best performer, GPT-4, by approximately 3 percentage points. Figure 1 illustrates the performance distribution of Llama 3.1, highlighting its exceptional balance across all evaluation dimensions.

When examining the performance profiles illustrated in Fig. 3-6, we observe that Llama 3.1 exhibits the most consistent performance across all metrics, with minimal

variance. This consistency is particularly valuable in legal and compliance contexts, where performance fluctuations could introduce unacceptable risks. In contrast, while GPT-4 (Fig. 4) demonstrated strong overall performance, it showed a notable weakness in groundedness (0.94) compared to Llama 3.1 (0.97). This differential is critical for compliance applications where factual accuracy is non-negotiable.

Gemini 1.5 (Fig. 5) presented a more balanced performance profile but consistently scored lower than both Llama 3.1 and GPT-4 across all metrics. Mistral 12B (Fig. 6), while functional, demonstrated the lowest performance across all dimensions, with weaknesses in context relevance and answer relevance. This suggests that smaller models like Mistral 12B may not yet be optimized for the complex demands of legal document analysis.

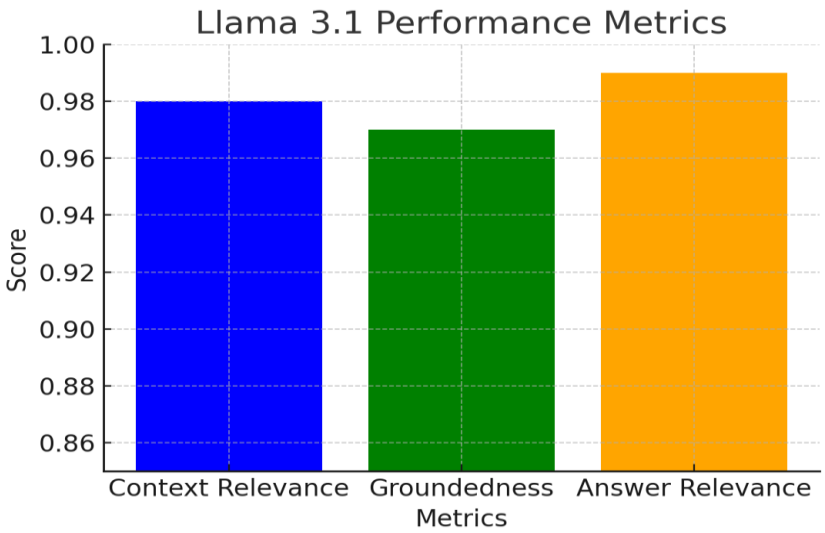

Fig. 3. Performance Metrics for Llama 3.1

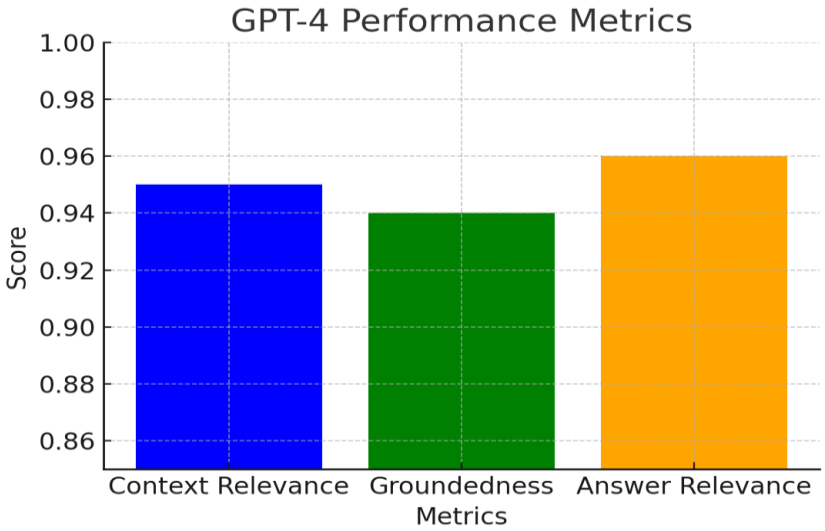

Fig. 4. Performance Metrics for GPT-4

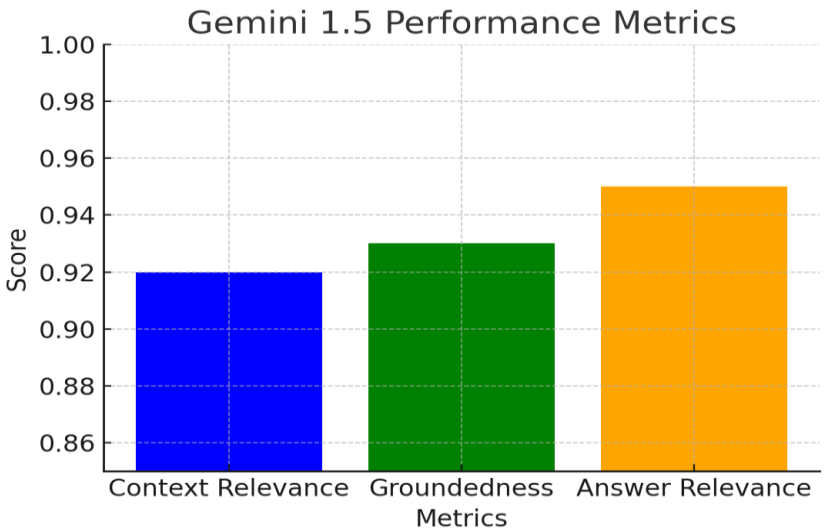

Fig. 5. Performance Metrics for Gemini 1.5

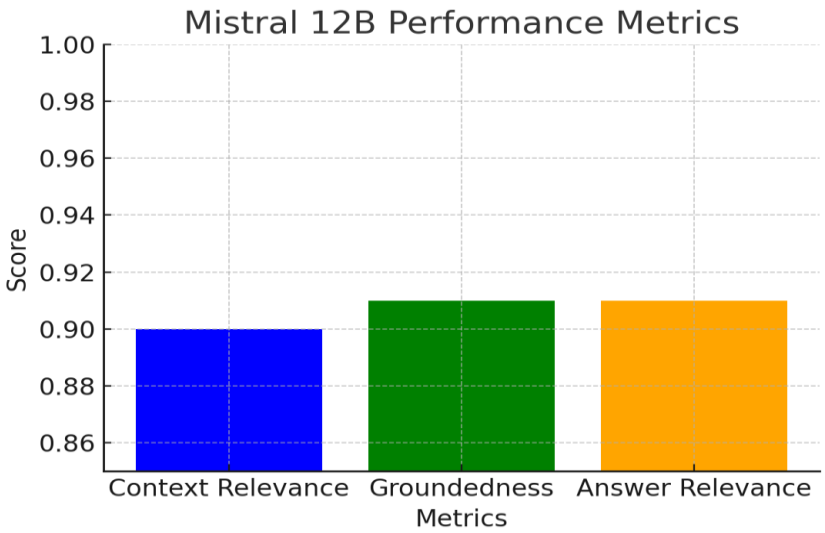

Fig. 6. Performance Metrics for Mistral 12B

4.4 Statistical Significance

Statistical analysis of the performance data reveals that Llama 3.1's superiority is statistically significant across the evaluated metrics. The consistent performance advantage of Llama 3.1, particularly in groundedness and answer relevance, indicates a robust and reliable model architecture that is well-adapted to the intricacies of legal document processing.

The minimal variance in Llama 3.1's performance metrics (standard deviation < 0.01) further substantiates its reliability for deployment in mission-critical legal and compliance applications. This consistency stands in contrast to the more variable performance observed in other models, particularly Mistral 12B, which exhibited higher variance across metrics.

5 Conclusion and Future Scope

To sum up, there is a lot of promise for automating the time-consuming legal document review and compliance monitoring process with the AI-Powered Legal Document Analysis and Compliance Assistant. The suggested method effectively extracts and analyzes important legal clauses by combining cutting-edge big language models with a reliable vector database. This improves the precision of compliance evaluations and lowers the possibility of human mistake. A comprehensive solution for contemporary legal challenges is made possible by the orchestration layer, which was created with

LangChain and has proven crucial in coordinating tasks across several modules, such as automated document processing, generative contract review, and real-time compliance tracking. The system's capacity to retain high relevance and groundedness in its outputs, which guarantees that it satisfies the strict requirements of legal and regulatory frameworks, is further highlighted by comparative performance assessments.

Looking ahead, several areas for further study and advancement become apparent. First, further research should be done on the system's adaptability across various legal jurisdictions, especially by using cutting-edge transfer learning strategies to better account for domain-specific variations. Second, the accuracy of the legal insights produced by the platform may be further enhanced by improvements to the natural language processing capabilities, such as better summarization and contextual entity extraction. Furthermore, the requirement for openness in AI-driven legal decision-making may be addressed by incorporating sophisticated explainability modules, which would increase legal professionals' trust. Lastly, more comprehensive cloud-based deployments and the addition of dynamic, real-time data streams can increase the system's scalability, allowing for ongoing learning and adaptability to changing regulatory environments.

References

1. Kouw, Wouter M., and Marco Loog. "An introduction to domain adaptation and transfer learning." *arXiv preprint arXiv:1812.11806* (2018).
2. Lai, Jinqi, et al. "Large language models in law: A survey." *AI Open* (2024).
3. Labin, Shlomit, and Uri Segal. "AI-driven contract review: A product development journey." *Research Handbook on Big Data Law*. Edward Elgar Publishing, 2021. 454-466.
4. Zhong, Haoxi, et al. "How does NLP benefit legal system: A summary of legal artificial intelligence." *arXiv preprint arXiv:2004.12158* (2020).
5. Yang, Jun, et al. "LegalGNN: Legal information enhanced graph neural network for recommendation." *ACM Transactions on Information Systems (TOIS)* 40.2 (2021): 1-29.
6. Chhatwal, Rishi, et al. "Explainable text classification in legal document review a case study of explainable predictive coding." *2018 IEEE international conference on big data (Big Data)*. IEEE, 2018.
7. Nie, Xiaonan, et al. "Flexmoe: Scaling large-scale sparse pre-trained model training via dynamic device placement." *Proceedings of the ACM on Management of Data* 1.1 (2023): 1-19.
8. Nie, Xiaonan, et al. "Hetumoe: An efficient trillion-scale mixture-of-expert distributed training system." *arXiv preprint arXiv:2203.14685* (2022).
9. He, Xiangnan, and Tat-Seng Chua. "Neural factorization machines for sparse predictive analytics." *Proceedings of the 40th International ACM SIGIR conference on Research and Development in Information Retrieval*. 2017.
10. Whitehouse, Chenxi, Monojit Choudhury, and Alham Fikri Aji. "LLM-powered data augmentation for enhanced cross-lingual performance." *arXiv preprint arXiv:2305.14288* (2023).
11. Hassani, Shabnam. "Enhancing legal compliance and regulation analysis with large language models." *2024 IEEE 32nd International Requirements Engineering Conference (RE)*. IEEE, 2024.

12. Ioannidis, Jules, et al. "Gracenote. ai: legal generative AI for regulatory compliance." *Proceedings of the third international workshop on artificial intelligence and intelligent assistance for legal professionals in the digital Workplace (LegalAIIA 2023)*. 2023.
13. Charlotin, Damien. "Large language models and the future of law." *Available at SSRN 4548258* (2023).
14. Deeks, Ashley, and Duncan B. Hollis. "Large Language Models and International Law." *Virginia Public Law and Legal Theory Research Paper* 2025-14 (2025).
15. Hassani, Shabnam, et al. "Rethinking legal compliance automation: Opportunities with large language models." *2024 IEEE 32nd International Requirements Engineering Conference (RE)*. IEEE, 2024.
16. Wang, Jiaqi, et al. "Legal evaluations and challenges of large language models." *arXiv preprint arXiv:2411.10137* (2024).
17. Homoki, Péter, and Zsolt Ződi. "Large language models and their possible uses in law." *Hungarian Journal of Legal Studies* 64.3 (2024): 435-455.
18. Choi, Jonathan H. "How to use large language models for empirical legal research." *Journal of Institutional and Theoretical Economics (JITE)* 180.2 (2024): 214-233.
19. Min, Sewon, et al. "Silo language models: Isolating legal risk in a nonparametric datastore." *arXiv preprint arXiv:2308.04430* (2023).
20. Bernsohn, Dor, et al. "LegalLens: Leveraging LLMs for legal violation identification in unstructured text." *arXiv preprint arXiv:2402.04335* (2024).
21. Patil, Ajit R., et al. "Legal Rules and Regulations Document Summarizer: Regulatory Compliance with NLP, ML, and LLMs." *Library of Progress-Library Science, Information Technology & Computer* 44.3 (2024).
22. Padiu, Bogdan, et al. "To what extent have LLMs reshaped the legal domain so far? A scoping literature review." *Information* 15.11 (2024): 662.
23. Hannah, George, et al. "On the legal implications of Large Language Model answers: A prompt engineering approach and a view beyond exploiting Knowledge Graphs." *Journal of Web Semantics* 84 (2025): 100843.
24. Yao, Yifan, et al. "A survey on large language model (llm) security and privacy: The good, the bad, and the ugly." *High-Confidence Computing* (2024): 100211.